

Brazilian E-Commerce Analytics

Order Fulfilment, Delivery Performance & Customer Satisfaction (2016–2018)

Sector	E-Commerce / Retail
Team	Section B • Group G-7
Faculty Mentor	Satyaki Das
Dataset	Olist Brazilian E-Commerce (Kaggle, 2016–2018)

Team Members

Saumya Kumar • Saman Iqbal Ansari • Arina Ali • Yash Kumar • Ayush Sahu • Prakhar Srivastava

GitHub: github.com/oksaumya/SectionB_G7_BrazilianE-Commerce

Tableau Public: <https://public.tableau.com/app/profile/saman.iqbal5804/viz/BrazilianEDA/SalesOverview?publish=yes&showOnboarding=true>

Context & Problem Statement

Why this matters and what decision the analysis supports

SECTOR CONTEXT

Brazilian e-commerce is growing fast on the back of smartphone penetration and expanding logistics but the country's geography and multi-vendor model create wide variance in delivery quality. Olist, Brazil's largest marketplace, sits at the centre of this tension between scale and reliability.

DECISION-MAKER

COO and Head of Seller Management of a large Brazilian online marketplace. They need to prioritise logistics investment, set seller onboarding standards, and decide where to spend customer-retention budget first.

CORE BUSINESS QUESTION

Which products, sellers, regions, and delivery patterns should management focus on to reduce service issues and improve revenue and repeat business?

OBJECTIVE

- Identify the operational drivers of late delivery, complaints, and revenue leakage.
- Quantify their impact and translate findings into 3–5 actionable recommendations the COO can fund.

Data Engineering — Source to Pipeline

From nine raw Olist tables to a single 100K-row analysis-ready master

SOURCE

- **Dataset:** Olist Brazilian E-Commerce (Kaggle)
- **Volume:** ~114,092 orders, 9 CSVs
- **Shape:** 100K+ rows × 37 cols (post-merge)
- **Period:** Sep 2016 to Aug 2018
- **Schema:** Relational, joined on order_id / customer_id / seller_id / product_id

PYTHON ETL (notebooks 01 & 02)

- **Nulls:** delivered_date 3.1% retained as NaN; review_score 0.8% excluded; category 0.1% → 'Unknown'
- **Outliers:** Negative delivery durations dropped; >\$5K orders kept (legit bulk)
- **Types:** All timestamps → datetime64; UF state codes standardised
- **Engineering:** delivery_days, delivery_delay_days, late_flag, order_month/quarter/year
- **Merge:** pd.merge() across 8 tables → olist_cleaned_master.csv

GITHUB & DICTIONARY


```
notebooks/  
  01_extraction.ipynb  
  02_cleaning.ipynb  
  03_eda.ipynb  
  04_statistical_analysis.ipynb  
  05_final_load_prep.ipynb  
  
data/  
  raw/      (9 CSVs)  
  processed/  
    cleaned_master.csv
```

KEY COLUMNS (analysis surface)

Column	Role	Column	Role
order_id	join key	review_score	satisfaction KPI
order_purchase_timestamp	time trends	payment_value	revenue
order_delivered_customer_date	delay calc	payment_type	payment behaviour
order_estimated_delivery_date	delivery gap	product_category_name	category perf.
order_status	fulfilment	seller_state / customer_state	geographic split

KPI & Metrics Framework

Five KPIs that map directly to the COO's decision levers



DEFINITIONS & FORMULAS

KPI	Why it matters	Formula	Notebook
Monthly Revenue Growth %	Tracks platform momentum	$(\text{Cur} - \text{Prev}) / \text{Prev} \times 100$	05_final_load_prep
Avg Customer Review Score	Customer satisfaction proxy	$\text{mean}(\text{review_score})$	03_edu
Late Delivery Rate	Fulfilment failure exposure	$\text{late_flag}=1 / \text{delivered} \times 100$	04_statistical
Cancellation Rate	Revenue leakage	$\text{cancelled} / \text{total} \times 100$	03_edu
On-Time Delivery Rate	Service-quality baseline	$\text{on_time} / \text{delivered} \times 100$	04_statistical

WHY THESE FIVE

Each KPI maps to one lever the COO actually controls growth, satisfaction, logistics, and revenue leakage — so every chart on the dashboard answers a question someone in the room can act on.

Key Insights from EDA

Six findings stated in business language — supporting charts in 03_eda.ipynb

1. Revenue grew ~4× year-on-year, but is concentrated.

Q1 2018 hit \$3.74M vs Q1 2017's \$0.94M. Top-3 categories (bed/bath/table, computers, health/beauty) drive a disproportionate share — disruption in any one materially impacts platform revenue.

2. São Paulo dominates both sides of the marketplace.

SP customers spend \$7.55M; SP sellers earn \$11.87M. Any uniform policy hits SP first and hardest — SP must be modelled as a separate operating regime.

3. Late deliveries hit satisfaction non-linearly past day 13.

Fast (<8d) reviews cluster above 4.0; slow (>13d) reviews fall below 3.0 with a steeper slope — each extra delay day costs more at higher delays.

4. Customers pay in instalments — credit card = 76.8% of revenue, ~3 instalments avg.

Revenue at order time ≠ cash realised; cash-flow planning for sellers depends on instalment lag, not order count.

5. Review distribution is bimodal, not normal.

5★ = 56.4%, but 1★ = 12.8% (13,194 orders) — vs only 3,484 at 2★. Dissatisfied customers report extreme dissatisfaction, consistent with delivery failures.

6. Q2 is the highest-revenue quarter every year.

\$3.84M in 2017 and comparable in 2018 — operations and seller readiness must be scaled into Q1 not after.

Advanced Analysis — Statistical Depth

Going beyond EDA: where is the delay-satisfaction relationship strongest, and who owns it?

DELIVERY TIME vs REVIEW SCORE — segmented regression

- **Fast (<8d):** review scores cluster above 4.0
- **Medium (8–13d):** scores 3.8–4.1, mixed
- **Slow (>13d):** scores below 3.0, steepest slope
- **Finding:** non-linear — each additional day costs more after the 13-day threshold

LATE-ORDER CONCENTRATION — by seller state

- **SP:** 6,730 late orders — 12× the next state (RJ 382)
- **MG / PR:** 481 / 550 late — secondary watch list
- **ES / BA:** 24 / 34 late — low absolute risk
- **Finding:** SP-specific intervention beats nationwide programmes on cost-of-impact

SEASONAL & CATEGORY EFFECTS

Late Order Volume by Month

Jan / Feb / Mar / Nov spikes (1,260 / 554 / 1,812 / 267); mid-year (Jun–Aug) shows lowest late counts. Avg delivery days drift 9.4 → 14.8 across the year — holiday-cycle pressure is visible.

Delivery Performance by Category

Audio, agro-industry, and art categories sit materially above the 12-day platform average — heavier or specialist logistics. Lightweight craft items deliver fastest.

Methods used: pandas + numpy aggregation, scipy/statsmodels for trend lines and segment comparisons, matplotlib/seaborn for visual validation — all in 04_statistical_analysis.ipynb.

Tableau Dashboard — Walkthrough

Three pages, six global filters, KPI-driven layout for executive and operational audiences

Page 1 — Sales Overview



Executive view: \$18.1M revenue, 103K orders, monthly trend, top categories.

Page 2 — Customers & Geography



86.88K unique customers, SP dominance, state-level revenue, review distribution.

Page 3 — Delivery & Operations



92.3% on-time, 8,756 late orders, delay-by-category, review vs delivery scatter.

GLOBAL FILTERS (live across all 3 pages): Year • State • Order Status • Payment Type • Product Category • Quarter

[TABLEAU PUBLIC URL](https://public.tableau.com/app/profile/saman.iqbal5804/viz/BrazilianEDA/SalesOverview?publish=yes&showOnboarding=true)

[<https://public.tableau.com/app/profile/saman.iqbal5804/viz/BrazilianEDA/SalesOverview?publish=yes&showOnboarding=true>]

Business Recommendations

Five prioritised actions each linked to an insight and an estimated impact

1

SP Seller Compliance Programme

Insight: Insight 5: SP sellers = 76.8% of late orders.

Action: Tiered scorecard; <90% on-time triggers review; <85% × 2 months → listing restrictions.

EXPECTED IMPACT

–673 late orders/period; +0.1–0.2 review-score points

2

Better Delivery Estimates + Proactive Comms

Insight: Insight 3: ~9-day average overrun on late orders.

Action: Re-fit estimate algorithm by state×category; auto-notify customers at 80% of window.

EXPECTED IMPACT

15–20% reduction in expectation-driven 1-star reviews

3

Top-3 Category Logistics SLA

Insight: Insight 1: top-3 categories drive disproportionate revenue & delays.

Action: Dedicated carrier SLAs + regional inventory pre-positioning for those categories.

EXPECTED IMPACT

\$200K–\$400K at-risk revenue protected per year

4

1-Star Customer Recovery Campaign

Insight: Insight 5b: 13,194 1-star reviewers — high churn risk.

Action: 48-hour personalised recovery: apology + voucher + expedited resolution; track 90-day repurchase.

EXPECTED IMPACT

~\$346K incremental revenue (15% recovery × \$175 AOV)

5

Pre-Q2 Demand Scaling

Insight: Insight 6: Q2 is peak, but Jan–Mar already spike late orders.

Action: Inventory readiness check 6 weeks pre-Q2; pre-negotiated carrier capacity for Jan–Mar.

EXPECTED IMPACT

≈350 fewer late orders per Q1–Q2; on-time rate → 93.7%

Impact & Value — The 'So What?'

What the platform stands to gain and what it loses by waiting

AGGREGATE IMPACT (low–high range, conservative assumptions)

\$546K+

Combined revenue at risk
protected or recovered annually

≈1,000

Fewer service failures
per Q1–Q2 cycle

+0.1–0.2

Points lift in average
review score

–15–20%

Reduction in expectation-
driven 1-star reviews

WHY ACT NOW

- **Compounding churn:** 1-star reviewers who don't return cost ~\$175 AOV × repeat-order multiplier each and Q1 alone has ~3,000 of them.
- **Concentration risk:** 76.8% of late orders sit with SP sellers; without intervention, every quarter's growth re-amplifies the same bottleneck.
- **Cheap wins are real:** Recommendation 2 (estimate accuracy) requires zero logistics investment only an algorithm and notification fix yet captures 15–20% of expectation-driven complaints.
- **Cost of inaction:** Even a 0.3-point review-score drop typically halves repeat-purchase probability; on \$18.1M base, the CLV exposure is in the multi-\$M range over 2 years.

Numbers are approximate; logic is documented in §11–12 of the project report.

Limitations & Next Steps

What we couldn't conclude and what should come next

LIMITATIONS

- **Correlation, not causation:** delivery vs review is observational; product quality, packaging, and seller comms are not controlled for.
- **Repeat-purchase weakness:** customer_id is per-order, not persistent robust repeat-rate computation requires deduplication via customer_unique_id.
- **Seller anonymisation:** no qualitative seller data (size, age, reputation) — we cannot diagnose why a seller underperforms.
- **Time truncation:** data ends Aug 2018 - no full-year 2018 view, no late-2018 holiday peak.
- **Order-level granularity:** no product-level margin or return rate unit economics out of scope.
- **Single platform:** findings apply to one marketplace; generalisation to Brazilian e-commerce broadly is not warranted.

NEXT STEPS

- **Predictive late-delivery model:** logistic regression / GBM on state × category × payment × timing for risk scoring at order placement.
- **Seller segmentation:** K-means on (on-time rate, review, revenue, volume) → tiered support programmes.
- **Review-text NLP:** sentiment + topic analysis on review comments to surface specific complaint drivers.
- **Real-time dashboard:** API-connected pipeline → operational alerting on emerging delay spikes.
- **Customer Lifetime Value:** link first-delivery experience to long-term repeat behaviour; quantify CLV loss per late delivery.

Team & Contribution

Visible commits, owned phases, declared in the report's contribution matrix

Member	Data	ETL	EDA	Stats	Tableau	Report	PPT
Saumya Kumar	Owner	—	Support	Support	—	—	Owner
Saman Iqbal Ansari	Support	Owner	Owner	Owner	Owner	—	—
Arina Ali	Support	—	—	—	—	Owner	—
Yash Kumar	—	—	—	Support	Support	—	—
Ayush Sahu	—	—	—	—	—	—	—
Prakhar Srivastava	—	—	—	—	—	—	—

PROJECT METADATA

Team Lead: Saumya Kumar
Mentor: Satyaki Das
Institute: Newton School of Technology
Section: B
Group: G-7
Date: April 2026

EVIDENCE TRAIL

- GitHub Insights: commit history per member
- Pull-request log: ownership of each phase
- Contribution matrix declared in §B of report
- All notebooks committed in .ipynb format

Thank You

Questions?

ONE-LINE TAKEAWAY

The Olist platform's growth story is real but operational reliability, concentrated in São Paulo sellers and a handful of categories, is the single highest-leverage decision surface for the COO over the next year.

GitHub: github.com/oksaumya/SectionB_G7_BrazilianE-Commerce

Tableau Public: <https://public.tableau.com/app/profile/saman.iqbal5804/viz/BrazillianEDA/SalesOverview?publish=yes&showOnboarding=true>